

In [1]:

```
import numpy as np
import pandas as pd
```

In [2]:

```
df = pd.read_csv('retail_price.csv')
```


In [3]:

```
df
```

Out[3]:

	product_id	product_category_name	month_year	qty	total_price	freight_price	unit_price	
0	bed1	bed_bath_table	01-05-2017	1	45.95	15.100000	45.950000	
1	bed1	bed_bath_table	01-06-2017	3	137.85	12.933333	45.950000	
2	bed1	bed_bath_table	01-07-2017	6	275.70	14.840000	45.950000	
3	bed1	bed_bath_table	01-08-2017	4	183.80	14.287500	45.950000	
4	bed1	bed_bath_table	01-09-2017	2	91.90	15.100000	45.950000	
...	
671	bed5	bed_bath_table	01-05-2017	1	215.00	8.760000	215.000000	
672	bed5	bed_bath_table	01-06-2017	10	2090.00	21.322000	209.000000	
673	bed5	bed_bath_table	01-07-2017	59	12095.00	22.195932	205.000000	
674	bed5	bed_bath_table	01-08-2017	52	10375.00	19.412885	199.509804	
675	bed5	bed_bath_table	01-09-2017	32	5222.36	24.324687	163.398710	

676 rows × 30 columns



Exploratory Data Analysis

In [4]:

```
df.shape
```

Out[4]:

(676, 30)

In [5]:

```
df.columns
```

Out[5]:

```
Index(['product_id', 'product_category_name', 'month_year', 'qty',  
      'total_price', 'freight_price', 'unit_price', 'product_name_lenght',  
      'product_description_lenght', 'product_photos_qty', 'product_weight_g',  
      'product_score', 'customers', 'weekday', 'weekend', 'holiday', 'month',  
      'year', 's', 'volume', 'comp_1', 'ps1', 'fp1', 'comp_2', 'ps2', 'fp2',  
      'comp_3', 'ps3', 'fp3', 'lag_price'],  
      dtype='object')
```

In [6]:

```
df.duplicated().sum()
```

Out[6]:

```
0
```

In [7]:

```
df.isnull().sum()
```

Out[7]:

```
product_id          0  
product_category_name  0  
month_year          0  
qty                0  
total_price         0  
freight_price       0  
unit_price          0  
product_name_lenght  0  
product_description_lenght  0  
product_photos_qty  0  
product_weight_g    0  
product_score       0  
customers           0  
weekday             0  
weekend             0  
holiday             0  
month               0  
year                0  
s                   0  
volume              0  
comp_1              0  
ps1                 0  
fp1                 0  
comp_2              0  
ps2                 0  
fp2                 0  
comp_3              0  
ps3                 0  
fp3                 0  
lag_price           0  
dtype: int64
```

In [8]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 676 entries, 0 to 675
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   product_id                           676 non-null    object
1   product_category_name                 676 non-null    object
2   month_year                           676 non-null    object
3   qty                                   676 non-null    int64
4   total_price                          676 non-null    float64
5   freight_price                        676 non-null    float64
6   unit_price                           676 non-null    float64
7   product_name_lenght                  676 non-null    int64
8   product_description_lenght           676 non-null    int64
9   product_photos_qty                   676 non-null    int64
10  product_weight_g                     676 non-null    int64
11  product_score                         676 non-null    float64
12  customers                             676 non-null    int64
13  weekday                              676 non-null    int64
14  weekend                                676 non-null    int64
15  holiday                              676 non-null    int64
16  month                                 676 non-null    int64
17  year                                  676 non-null    int64
18  s                                     676 non-null    float64
19  volume                               676 non-null    int64
20  comp_1                              676 non-null    float64
21  ps1                                  676 non-null    float64
22  fp1                                  676 non-null    float64
23  comp_2                              676 non-null    float64
24  ps2                                  676 non-null    float64
25  fp2                                  676 non-null    float64
26  comp_3                              676 non-null    float64
27  ps3                                  676 non-null    float64
28  fp3                                  676 non-null    float64
29  lag_price                            676 non-null    float64
dtypes: float64(15), int64(12), object(3)
memory usage: 158.6+ KB
```

In [9]:

```
df.describe()
```

Out[9]:

	qty	total_price	freight_price	unit_price	product_name_lenght	product_descript
count	676.000000	676.000000	676.000000	676.000000	676.000000	
mean	14.495562	1422.708728	20.682270	106.496800	48.720414	
std	15.443421	1700.123100	10.081817	76.182972	9.420715	
min	1.000000	19.900000	0.000000	19.900000	29.000000	
25%	4.000000	333.700000	14.761912	53.900000	40.000000	
50%	10.000000	807.890000	17.518472	89.900000	51.000000	
75%	18.000000	1887.322500	22.713558	129.990000	57.000000	
max	122.000000	12095.000000	79.760000	364.000000	60.000000	3i

8 rows × 27 columns



In [10]:

```
df.nunique()
```

Out[10]:

product_id	52
product_category_name	9
month_year	20
qty	66
total_price	573
freight_price	653
unit_price	280
product_name_lenght	24
product_description_lenght	46
product_photos_qty	7
product_weight_g	45
product_score	11
customers	94
weekday	4
weekend	3
holiday	5
month	12
year	2
s	450
volume	40
comp_1	88
ps1	9
fp1	179
comp_2	123
ps2	10
fp2	242
comp_3	105
ps3	9
fp3	229
lag_price	307
dtype: int64	

In [11]:

```
object_columns = df.select_dtypes(include='object').columns
print("Object Columns:")
print(object_columns)
print()

# Identify numerical types
numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns
print("Numerical Columns:")
print(numerical_columns)
```

Object Columns:
Index(['product_id', 'product_category_name', 'month_year'], dtype='object')

Numerical Columns:
Index(['qty', 'total_price', 'freight_price', 'unit_price',
 'product_name_lenght', 'product_description_lenght',
 'product_photos_qty', 'product_weight_g', 'product_score', 'customers',
 'weekday', 'weekend', 'holiday', 'month', 'year', 's', 'volume',
 'comp_1', 'ps1', 'fp1', 'comp_2', 'ps2', 'fp2', 'comp_3', 'ps3', 'fp3',
 'lag_price'],
 dtype='object')

In [12]:

```
import matplotlib.pyplot as plt
import seaborn as sns
```

In [13]:

```
import warnings
warnings.filterwarnings('ignore')
```

In [14]:

```
for i in object_columns:
    print(i,':')
    print(df[i].unique())
    print('\n')
```

```
product_id :
['bed1' 'garden5' 'consoles1' 'garden7' 'health9' 'cool4' 'health3'
 'perfumery1' 'cool5' 'health8' 'garden4' 'computers5' 'garden10'
 'computers6' 'health6' 'garden6' 'health10' 'watches2' 'health1'
 'garden8' 'garden9' 'watches6' 'cool3' 'perfumery2' 'cool2' 'computers1'
 'consoles2' 'health5' 'watches8' 'furniture4' 'watches5' 'health7' 'bed3'
 'garden3' 'bed2' 'furniture3' 'watches4' 'watches3' 'furniture2'
 'garden2' 'furniture1' 'health2' 'garden1' 'cool1' 'computers4'
 'watches7' 'computers3' 'health4' 'watches1' 'computers2' 'bed4' 'bed5']
```

```
product_category_name :
['bed_bath_table' 'garden_tools' 'consoles_games' 'health_beauty'
 'cool_stuff' 'perfumery' 'computers_accessories' 'watches_gifts'
 'furniture_decor']
```

```
month_year :
['01-05-2017' '01-06-2017' '01-07-2017' '01-08-2017' '01-09-2017'
 '01-10-2017' '01-11-2017' '01-12-2017' '01-01-2018' '01-02-2018'
 '01-03-2018' '01-04-2018' '01-05-2018' '01-06-2018' '01-07-2018'
 '01-08-2018' '01-03-2017' '01-04-2017' '01-02-2017' '01-01-2017']
```

In [15]:

```
for i in object_columns:  
    print(i,':')  
    print(df[i].value_counts())  
    print('\n')
```



```
product_id :
health5      20
health7      20
bed2         19
garden1      18
health9      18
garden3      18
computers4   18
health8      17
watches1     17
garden9      17
garden2      17
garden7      16
garden10     16
garden6      16
bed1         16
computers1   15
cool1        15
watches3     15
watches2     15
garden5      14
garden4      14
garden8      14
watches6     14
perfumery2   13
cool2        13
furniture2   13
health2      13
furniture1   13
perfumery1   13
cool5        13
watches7     12
furniture3   12
consoles1    12
health4      11
bed3         11
computers3   10
computers2   10
bed4         10
consoles2    10
watches4     10
watches5     10
furniture4   10
watches8     10
health1      9
cool4        9
computers6   8
computers5   8
health3      8
cool3        7
health10     7
health6      7
bed5         5
Name: product_id, dtype: int64
```

```
product_category_name :
garden_tools      160
health_beauty     130
watches_gifts     103
computers_accessories  69
bed_bath_table    61
```


```
cool_stuff          57
furniture_decor     48
perfumery           26
consoles_games      22
Name: product_category_name, dtype: int64
```

```
month_year :
01-03-2018    50
01-02-2018    49
01-01-2018    48
01-04-2018    48
01-11-2017    44
01-12-2017    44
01-10-2017    43
01-06-2018    42
01-05-2018    40
01-07-2018    40
01-08-2018    38
01-08-2017    37
01-09-2017    36
01-07-2017    33
01-06-2017    25
01-05-2017    20
01-04-2017    15
01-03-2017    13
01-02-2017     9
01-01-2017     2
Name: month_year, dtype: int64
```


In [16]:

```
for i in object_columns:
    print(i,':')
    plt.figure(figsize=(15,6))
    sns.countplot(df[i], data = df, palette = 'hls')
    plt.xticks(rotation = 90)
    plt.show()
    print('\n')
```


product_id :




product_category_name :



month_year :



month_year :



In [18]:


```
import plotly.express as px
import plotly.graph_objects as go
```

In [19]:

```
for i in object_columns:
    print(i, ':')
    fig = go.Figure(data=[go.Bar(x=df[i].value_counts().index, y=df[i].value_counts())])
    fig.update_layout(
        title=i,
        xaxis_title="Categories",
        yaxis_title="Count"
    )
    fig.show()
    print('\n')
```

product_id :


product_id



In [20]:


```
for i in object_columns:
    print(i, ':')
    counts = df[i].value_counts()
    fig = go.Figure(data=[go.Pie(labels=counts.index, values=counts)])
    fig.update_layout(title=i)
    fig.show()
    print('\n')
```

```
product_id :
```




In [21]:

```
for i in numerical_columns:
    plt.figure(figsize=(15,6))
    sns.histplot(df[i], kde = True, bins = 20, palette = 'hls')
    plt.xticks(rotation = 90)
    plt.show()
```




In [22]:

```
for i in numerical_columns:  
    plt.figure(figsize=(15,6))  
    sns.distplot(df[i], kde = True, bins = 20)  
    plt.xticks(rotation = 90)  
    plt.show()
```




In [23]:

```
for i in numerical_columns:  
    plt.figure(figsize=(15,6))  
    sns.boxplot(df[i], data = df, palette = 'hls')  
    plt.xticks(rotation = 90)  
    plt.show()
```



In [24]:

```
for i in numerical_columns:
    plt.figure(figsize=(15,6))
    sns.violinplot(df[i], data = df, palette = 'hls')
    plt.xticks(rotation = 90)
    plt.show()
```

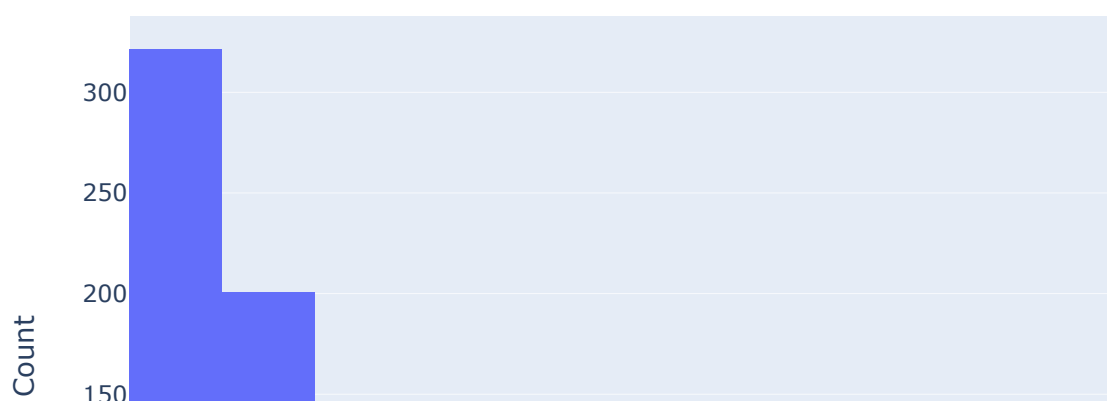


In [25]:

```
for i in numerical_columns:
    print(i, ':')
    fig = go.Figure(data=[go.Histogram(x=df[i], nbinsx=20)])
    fig.update_layout(
        title=i,
        xaxis_title="Value",
        yaxis_title="Count"
    )
    fig.show()
    print('\n')
```

qty :

qty




In [26]:

```
for i in numerical_columns:
    print(i, ':')
    fig = go.Figure(data=[go.Box(x=df[i])])
    fig.update_layout(
        title=i + ' Box Plot',
        xaxis_title="Value"
    )
    fig.show()
    print('\n')
```

qty :

qty Box Plot



In [27]:


```

for i in numerical_columns:
    print(i, ':')
    fig = go.Figure(data=[go.Violin(y=df[i])])
    fig.update_layout(
        title=i + ' Violin Plot',
        yaxis_title="Value"
    )
    fig.show()
    print('\n')

```

qty :

qty Violin Plot





In [28]:

```


for i in numerical_columns:
    for j in object_columns:
        plt.figure(figsize=(15,6))
        sns.barplot(x = df[j], y = df[i], data = df, ci = None, palette = 'hls')
        plt.xticks(rotation = 90)
        plt.show()

```



In [29]:

```
for i in numerical_columns:
    for j in object_columns:
        plt.figure(figsize=(15,6))
        sns.boxplot(x = df[j], y = df[i], data = df, palette = 'hls')
        plt.xticks(rotation = 90)
        plt.show()
```




In [30]:

```
for i in numerical_columns:
    # Iterate over the object columns
    for j in object_columns:
        # Create a bar plot using Plotly
        fig = px.bar(df, x=j, y=i)

        # Set the figure size
        fig.update_layout(width=800, height=400)

        # Rotate x-axis labels if necessary
        fig.update_layout(xaxis_tickangle=-45)

        # Show the plot
        fig.show()
```




In [31]:

```
for i in numerical_columns:
    # Iterate over the object columns
    for j in object_columns:
        # Create a box plot using Plotly
        fig = px.box(df, x=j, y=i)

        # Set the figure size
        fig.update_layout(width=800, height=400)


        # Rotate x-axis labels if necessary
        fig.update_layout(xaxis_tickangle=-45)

        # Show the plot
        fig.show()
```




In [32]:

```
for i in numerical_columns:
    for j in numerical_columns:
        if i != j:
            plt.figure(figsize=(15,6))
            sns.lineplot(x = df[j], y = df[i], data = df, ci = None, palette = 'hls')
            plt.xticks(rotation = 90)
            plt.show()
```



In [33]:

```
for i in numerical_columns:
    for j in numerical_columns:
        if i != j:
            plt.figure(figsize=(15,6))
            sns.scatterplot(x = df[j], y = df[i], data = df, ci = None, palette = 'hls')
            plt.xticks(rotation = 90)
            plt.show()
```



Calculate Revenue and Profit:

In [34]:

```
df['revenue'] = df['qty'] * df['total_price']
df['profit'] = df['revenue'] - df['freight_price']
```

Calculate Margin:

In [35]:

```
df['margin'] = (df['profit'] / df['revenue']) * 100
```

Price Ratios:

In [36]:

```
df['price_ratio_1'] = df['unit_price'] / df['comp_1']
df['price_ratio_2'] = df['unit_price'] / df['comp_2']
df['price_ratio_3'] = df['unit_price'] / df['comp_3']
```

Price Differences:

In [37]:

```
df['price_diff_1'] = df['unit_price'] - df['comp_1']
df['price_diff_2'] = df['unit_price'] - df['comp_2']
df['price_diff_3'] = df['unit_price'] - df['comp_3']
```


In [38]:

```
df.head()
```

Out[38]:

	product_id	product_category_name	month_year	qty	total_price	freight_price	unit_price	profit
0	bed1	bed_bath_table	01-05-2017	1	45.95	15.100000	45.95	
1	bed1	bed_bath_table	01-06-2017	3	137.85	12.933333	45.95	
2	bed1	bed_bath_table	01-07-2017	6	275.70	14.840000	45.95	
3	bed1	bed_bath_table	01-08-2017	4	183.80	14.287500	45.95	
4	bed1	bed_bath_table	01-09-2017	2	91.90	15.100000	45.95	

5 rows × 39 columns



Market Demand Indicators:

In [39]:

```
df['customer_score_ratio'] = df['customers'] / df['product_score']
df['customer_photo_ratio'] = df['customers'] / df['product_photos_qty']
df['description_length_ratio'] = df['product_description_lenght'] / df['product_name_lenght']
```

In [40]:

```
df.head()
```

Out[40]:

	product_id	product_category_name	month_year	qty	total_price	freight_price	unit_price	pro
0	bed1	bed_bath_table	01-05-2017	1	45.95	15.100000	45.95	
1	bed1	bed_bath_table	01-06-2017	3	137.85	12.933333	45.95	
2	bed1	bed_bath_table	01-07-2017	6	275.70	14.840000	45.95	
3	bed1	bed_bath_table	01-08-2017	4	183.80	14.287500	45.95	
4	bed1	bed_bath_table	01-09-2017	2	91.90	15.100000	45.95	

5 rows × 42 columns

Time-related Features:

In [41]:

```
df['month_year'] = pd.to_datetime(df['month_year'])
df['month'] = df['month_year'].dt.month
df['year'] = df['month_year'].dt.year
df['is_weekend'] = df['weekday'].apply(lambda x: 1 if x >= 5 else 0)
df['is_holiday'] = df['holiday']
```

In [42]:

```
df.head()
```

Out[42]:

	product_id	product_category_name	month_year	qty	total_price	freight_price	unit_price	pro
0	bed1	bed_bath_table	2017-01-05	1	45.95	15.100000	45.95	
1	bed1	bed_bath_table	2017-01-06	3	137.85	12.933333	45.95	
2	bed1	bed_bath_table	2017-01-07	6	275.70	14.840000	45.95	
3	bed1	bed_bath_table	2017-01-08	4	183.80	14.287500	45.95	
4	bed1	bed_bath_table	2017-01-09	2	91.90	15.100000	45.95	

5 rows × 44 columns

Lagged Price:

In [43]:

```
df['lag_price'] = df.groupby('product_id')['total_price'].shift(1)
```

Handling Categorical Variables:

In [44]:

```
# One-hot encoding example
df_encoded = pd.get_dummies(df, columns=['product_category_name', 'weekday'])
```

Scaling Numeric Features:

In [45]:

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
numeric_features = ['qty', 'total_price', 'freight_price', 'unit_price', 'product_name_length',
                    'product_description_length', 'product_photos_qty', 'product_weight_g',
                    'product_score', 'customers', 'volume', 'comp_1', 'ps1', 'fp1', 'comp_2',
                    'ps2', 'fp2', 'comp_3', 'ps3', 'fp3', 'lag_price']

df_scaled = df_encoded.copy()
df_scaled[numeric_features] = scaler.fit_transform(df_encoded[numeric_features])
```

Splitting the Dataset:

In [46]:

```

from sklearn.model_selection import train_test_split
exclude_columns = ['month_year']
X = df_scaled.drop(['total_price', 'product_id'] + exclude_columns, axis=1)
y = df_scaled['total_price']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

In [47]:

```

subset_cols = ['qty', 'total_price', 'freight_price', 'unit_price', 'product_score', 'customers']
subset_df = df[subset_cols]

corr_matrix = subset_df.corr()

plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()


sns.pairplot(subset_df, vars=['qty', 'total_price', 'unit_price', 'product_score', 'customers'])
plt.show()

plt.figure(figsize=(12, 8))
sns.boxplot(x='product_category_name', y='total_price', data=df)
plt.title('Product Category vs. Total Price')
plt.xticks(rotation=90)
plt.show()

plt.figure(figsize=(10, 6))
sns.barplot(x='weekday', y='total_price', data=df)
plt.title('Weekday vs. Total Price')
plt.show()

plt.figure(figsize=(8, 6))
sns.countplot(x='holiday', data=df, hue='total_price')
plt.title('Holiday vs. Total Price')
plt.show()

```




In [48]:

```
plt.figure(figsize=(12, 8))
df.groupby('product_category_name')['total_price'].mean().sort_values().plot(kind='bar')
plt.title('Product Category vs. Average Total Price')
plt.xlabel('Product Category')
plt.ylabel('Average Total Price')
plt.xticks(rotation=90)
plt.show()

plt.figure(figsize=(12, 8))
df.groupby('month_year')['total_price'].sum().plot(kind='line')
plt.title('Total Price Trend over Time')
plt.xlabel('Month-Year')
plt.ylabel('Total Price')
plt.xticks(rotation=45)
plt.show()

plt.figure(figsize=(10, 6))
plt.hist(df['unit_price'], bins=20, edgecolor='k')
plt.title('Distribution of Unit Price')
plt.xlabel('Unit Price')
plt.ylabel('Frequency')
plt.show()

plt.figure(figsize=(10, 6))
plt.scatter(df['product_weight_g'], df['total_price'])
plt.title('Product Weight vs. Total Price')
plt.xlabel('Product Weight (grams)')
plt.ylabel('Total Price')
plt.show()
```



In [49]:

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

model = LinearRegression()
```

In [50]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

In [51]:

```
from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='mean')
```

In [52]:

```
X_train_imputed = imputer.fit_transform(X_train)
X_test_imputed = imputer.transform(X_test)
```

In [53]:

```
model.fit(X_train_imputed, y_train)
```

Out[53]:

```
LinearRegression
LinearRegression()
```

In [54]:

```
y_pred = model.predict(X_test_imputed)
```

In [55]:

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Calculate mean squared error (MSE)
mse = mean_squared_error(y_test, y_pred)

# Calculate root mean squared error (RMSE)
rmse = mean_squared_error(y_test, y_pred, squared=False)


# Calculate R-squared score
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error (MSE):", mse)
print("Root Mean Squared Error (RMSE):", rmse)
print("R-squared Score:", r2)
```

```
Mean Squared Error (MSE): 0.17492913184386694
Root Mean Squared Error (RMSE): 0.41824530104218377
R-squared Score: 0.8401763893992591
```

In [56]:

```
# Plotting the predicted values against the actual values  
plt.scatter(y_test, y_pred)  
plt.xlabel('Actual Values')  
plt.ylabel('Predicted Values')  
plt.title('Actual vs Predicted Values')  
plt.show()
```



In [57]:


```
# Create a DataFrame with actual and predicted values
df_plot = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})

# Create a scatter plot
fig = px.scatter(df_plot, x='Actual', y='Predicted', title='Actual vs Predicted Values')

# Add a diagonal line for reference
fig.add_shape(type='line', x0=df_plot['Actual'].min(), y0=df_plot['Actual'].min(),
              x1=df_plot['Actual'].max(), y1=df_plot['Actual'].max(),
              line=dict(color='red', dash='dash'))

# Show the plot
fig.show()
```

Actual vs Predicted Values



In [58]:

```
from sklearn.tree import DecisionTreeRegressor

model = DecisionTreeRegressor()
```


In [59]:

```
model.fit(X_train_imputed, y_train)
```

Out[59]:

```
▼ DecisionTreeRegressor  
DecisionTreeRegressor()
```

In [60]:

```
y_pred = model.predict(X_test_imputed)
```


In [61]:

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score  
  
# Calculate mean squared error (MSE)  
mse = mean_squared_error(y_test, y_pred)  
  
# Calculate root mean squared error (RMSE)  
rmse = mean_squared_error(y_test, y_pred, squared=False)  
  
# Calculate R-squared score  
r2 = r2_score(y_test, y_pred)  
  
print("Mean Squared Error (MSE):", mse)  
print("Root Mean Squared Error (RMSE):", rmse)  
print("R-squared Score:", r2)
```

```
Mean Squared Error (MSE): 0.06924014365516341  
Root Mean Squared Error (RMSE): 0.26313521933630135  
R-squared Score: 0.9367388973989801
```

In [62]:

```
# Plotting the predicted values against the actual values  
plt.scatter(y_test, y_pred)  
plt.xlabel('Actual Values')  
plt.ylabel('Predicted Values')  
plt.title('Actual vs Predicted Values')  
plt.show()
```



In [63]:

```
# Create a DataFrame with actual and predicted values
df_plot = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})

# Create a scatter plot
fig = px.scatter(df_plot, x='Actual', y='Predicted', title='Actual vs Predicted Values')

# Add a diagonal line for reference
fig.add_shape(type='line', x0=df_plot['Actual'].min(), y0=df_plot['Actual'].min(),
              x1=df_plot['Actual'].max(), y1=df_plot['Actual'].max(),
              line=dict(color='red', dash='dash'))

# Show the plot
fig.show()
```

Actual vs Predicted Values

